

PeFoMed-Light: A Lightweight and Parameter-Efficient Vision–Language Architecture for Automated Medical Report Generation

Medical Report Generation using PeFoMed with LLM-based Text Reports from CheXpert Dataset

Presenters:

Ch. Himesh Raj Kumar (CS23BT005)

K. Varaprasad Reddy (CS23BT050)

Pasham Pardeev Narsi Reddy(CS22BT040)

A research project on improving automated medical report generation using parameter-efficient fine-tuning and large-scale, region-based datasets.



Deep Learning



Medical AI



Report Generation

CONTENTS

- Abstract
- Introduction
- Methodology
- Experiments and Results
- Evaluation Metrics
- Conclusion and Future work
- References

Abstract

- The project focuses on automating medical report generation from chest X-ray images using the PeFoMed model.
- PeFoMed was fine-tuned on a custom dataset derived from CheXpert.
- Llama was used to generate text reports corresponding to the 14 diagnostic labels in the dataset.
- The model was trained with:
 - Chest X-ray images as input, and
 - LLM-generated text reports as output.
- The main goal was to evaluate PeFoMed's ability to generate accurate medical reports from image inputs.

Introduction

- PeFoMed is a **transformer-based multimodal model** designed for **medical report generation**.
- Traditional medical report generation models are **limited by high data annotation costs**.
- To overcome this, **LLM-generated synthetic text reports** aligned with **CheXpert labels** are used to **create a scalable training dataset** for PeFoMed.
- The project aims to:
 - **Verify** whether LLM-generated medical text **accurately aligns with ground truth reports**.
 - **Evaluate** model performance using **standard NLP metrics**.

What is PeFoMed? and Why PeFoMed?

- PeFoMed is a transformer-based multimodal generative model designed for medical imaging tasks like Visual Question Answering (VQA) and Medical Report Generation (MRG).
It follows a parameter-efficient fine-tuning (PEFT) strategy so that only a small part of the model is trained, while the rest stays frozen.
- **Architecture Overview-**
- **Vision Encoder:** A pre-trained ViT (EVA) model extracts features from X-ray images (kept *frozen*).
- **Projection Layer:** Maps visual features into the LLM embedding space.
- **LLM Backbone:** A frozen LLaMA-2-Chat model, acting as a text generator.
- **LoRA Layers:** Lightweight trainable adapters inserted into the LLM to enable efficient fine-tuning.
- **Why PeFoMed?**
- Requires very little training compute
- Maintains strong generative capability
- Works well even with limited medical datasets
- Proven effective on VQA-RAD, SLAKE, PathVQA, IU-Xray

Text Generation Pipeline

- We generate Findings section from CheXpert labels using LangChain + LLaMA-7b with few-shot prompting
 - **Input:** CheXpert label vector (14 labels)
 - **Output:** Radiology-style Findings text.
- **Pipeline Flow:**
- 1.Load CheXpert labels
 - 2.Convert them into a natural language prompt
 - 3.LLaMA-7b generates the Findings section
 - 4.Output checked manually + validated using CheXbert
- **Advantages:**
- Zero hallucination observed in manual check
 - Radiologist-like tone preserved
 - Works consistently across normal and abnormal cases

Why We Modified the Original PeFoMed Architecture

Problem in Original PeFoMed:

- The PeFoMed paper uses EVA-ViT-g Vision Encoder (a very large ViT model), its not publicly available and requires huge GPU resources.
- EVA produces visual embeddings that are already aligned to 5632 dimensions, which match LLaMA hidden size requirements.
- The original PeFoMed uses a large frozen LLaMA-Chat model for text decoding, very high VRAM (>40GB).

Resolvment of Issue in Our Implementation:

- EVA-ViT is not open-source, so we replaced it with: CLIP ViT-L/14 Vision Encoder (OpenAI)
- But CLIP produces: token embedding size = 1024, We take 4 tokens : $4 \times 1024 = 4096$ dimensions(5632 for EVA-ViT)
- Our Solution: To replicate the original structure: We added an extra trainable projection layer:
- CLIP output (4096)->->-(up_proj)->->->5632->->->(image_proj)->->->4096(LLaMA hidden embedding size)

Contd..

- But, Due to limited hardware resources, we were unable to load or fine-tune the complete LLaMA decoder so used a smaller text decoder(GPT-2)

GPT-2 lightweight decoder generates reports:

- Fully compatible with our low-resource hardware
- Entire pipeline trained **from scratch for our dataset**
- Base GPT-2 model weights remain **frozen**
- Only **LoRA adapters** are trained (parameter-efficient fine-tuning)

Final Result

- We successfully maintain PeFoMed architecture logic
- But with a fully open-source vision encoder (CLIP) and open source text decoder (GPT-2)
- And minimal trainable parameters (up_proj+image_proj+LoRA)

Final Model Pipeline (Modified: CLIP + Projection + GPT-2 Text Decoder)

1. Image Processing

- Input image goes into CLIP ViT-L/14 (frozen).
- Outputs 4096-dim visual embedding.

2. Visual Feature Alignment

- Up-projection layer: $4096 \rightarrow 5632$ (trainable).
- Image-projection layer: $5632 \rightarrow 4096$ (trainable).

3. Multimodal Fusion

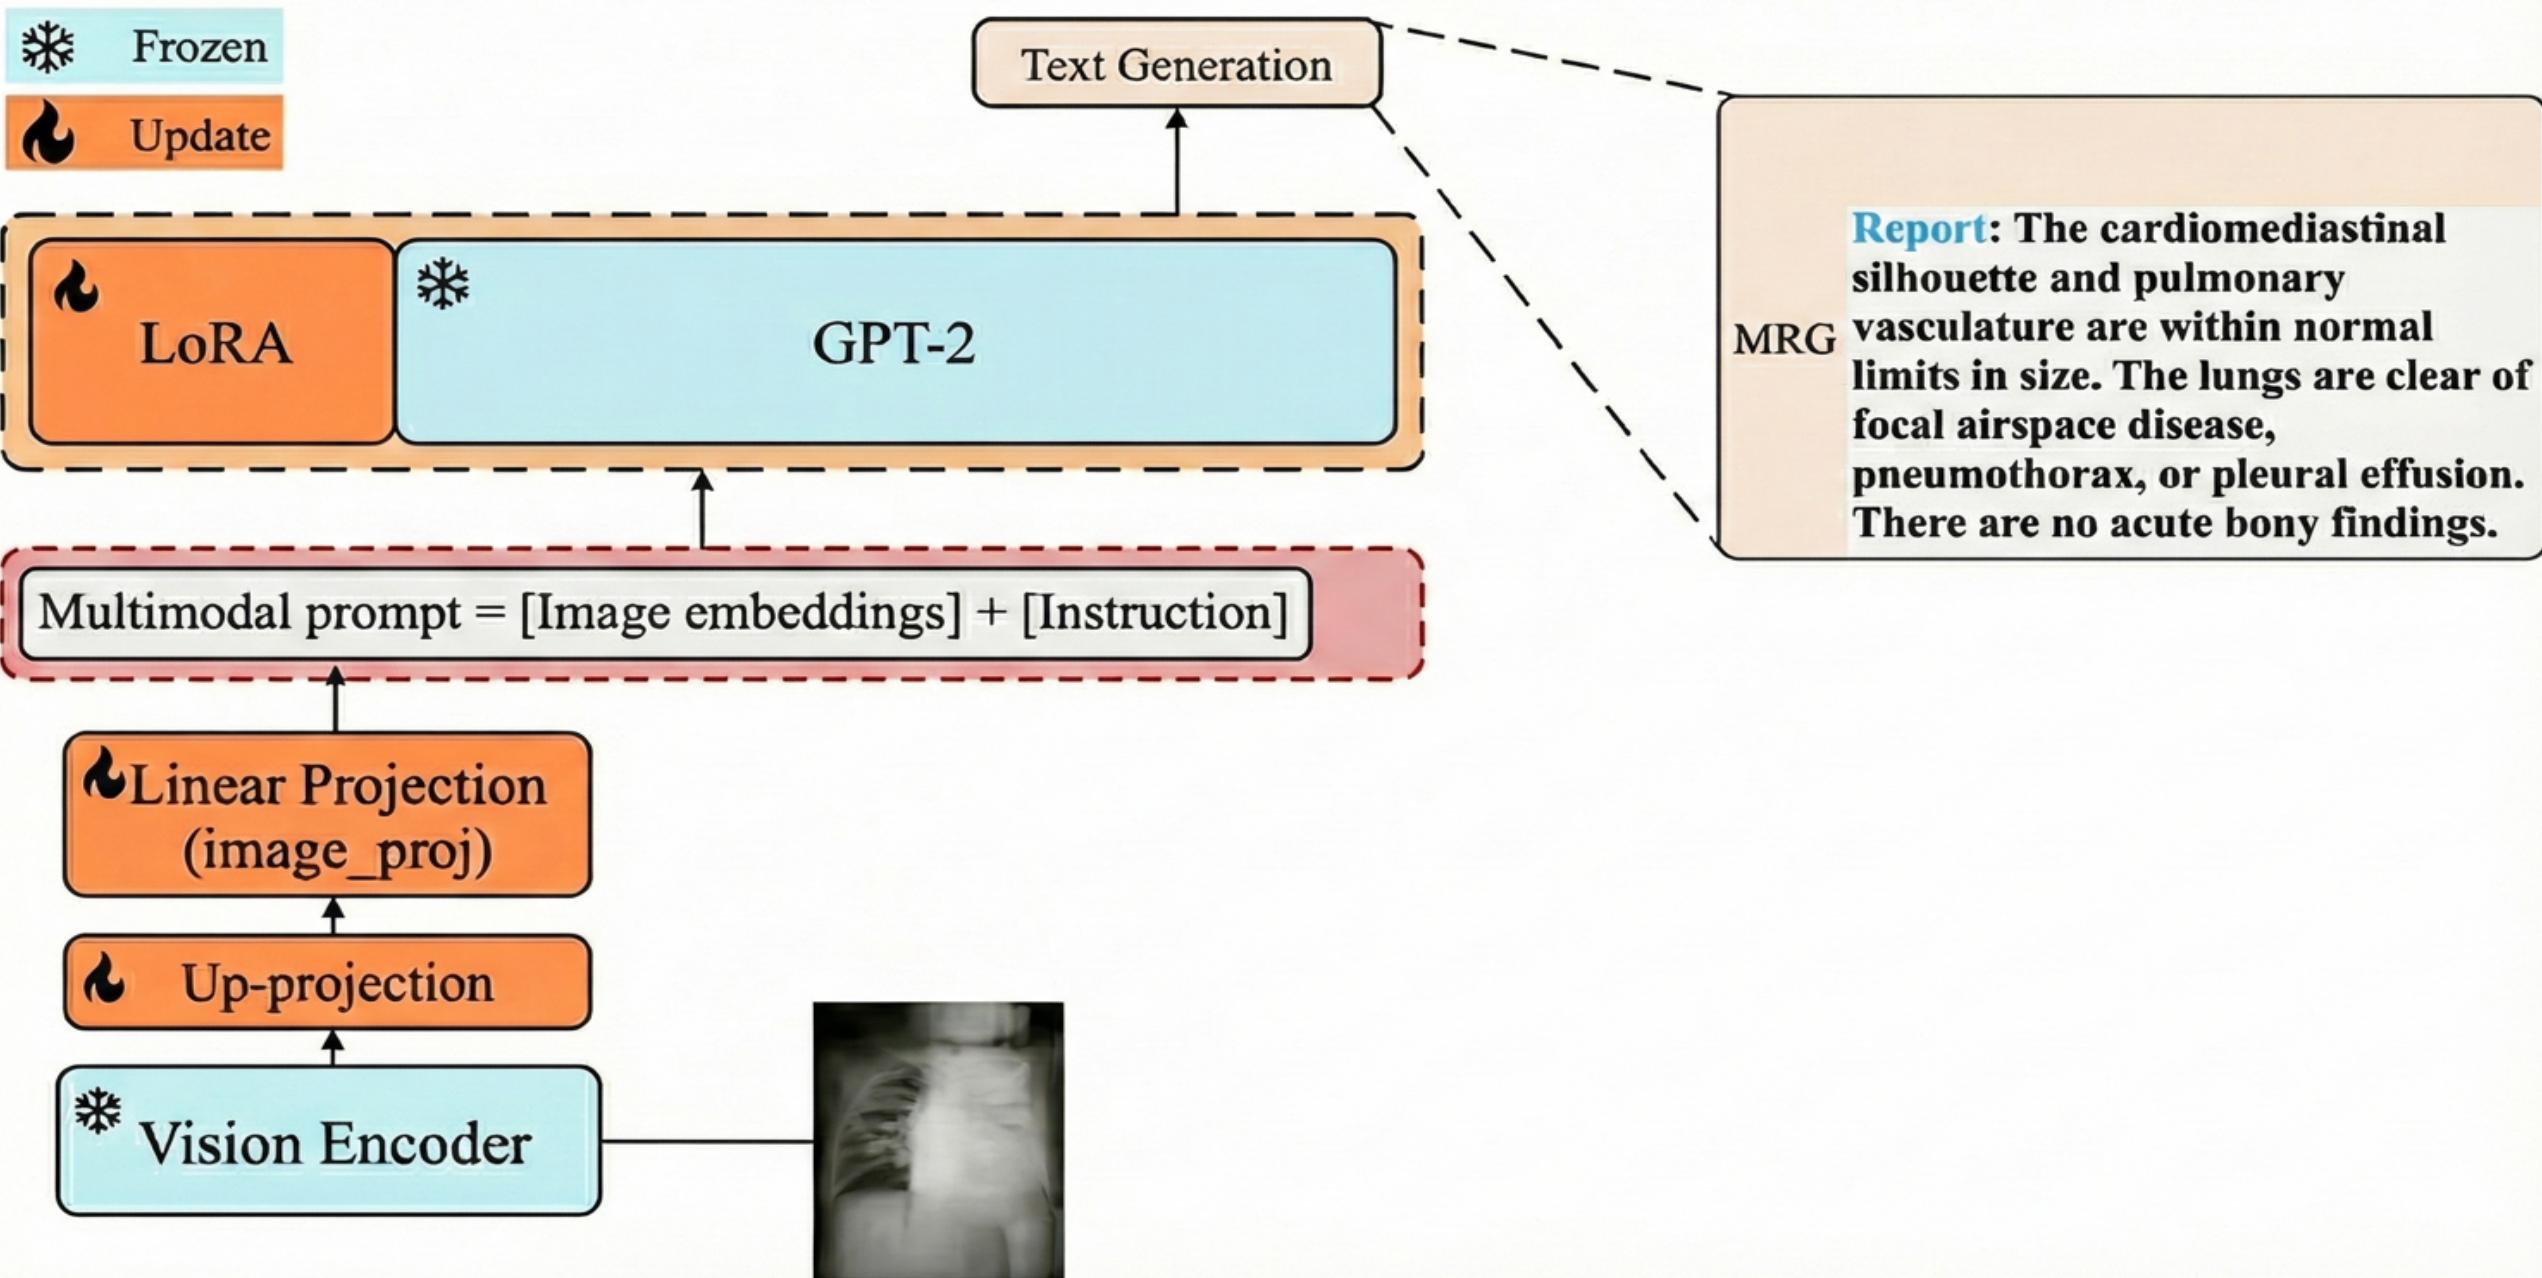
- Aligned visual tokens are injected into the GPT-2 decoder.

4. Text Generation

- GPT-2 generates answers/reports.
- Base GPT-2 weights frozen.
- Only LoRA adapters are trained.

5. Efficiency

- Only projection layers + LoRA are trainable.
- Achieves parameter-efficient fine-tuning.



Mathematics Behind PeFoMed

- **Visual Embedding (ViT EVA):** Chest X-ray image → split into patches → patch embeddings E_{img}

$$E_{img} = \text{ViT}(X)$$

- **Token Grouping:** 4 consecutive visual tokens are concatenated to reduce computation.

$$G_i = \text{Concat}(E_{4i}, E_{4i+1}, E_{4i+2}, E_{4i+3})$$

- **Projection into LLM Space:** A single linear layer maps visual features into GPT-2 hidden space.

$$Z_i = W_p G_i + b_p$$

- **LoRA Update Rule (Parameter-Efficient Fine-Tuning):** Instead of updating the full weight W , LoRA learns low-rank matrices-

$$W' = W + BA, \quad \text{where } A \in \mathbb{R}^{r \times d}, B \in \mathbb{R}^{d \times r}$$

- **Prompt Fusion:** $\text{M}^{\text{multimodal prompt}} = [\text{Image embedding}] \cdot [\text{Task token}] \cdot [\text{Instruction}]$

$$T = [Z_1, Z_2, \dots, Z_n, [\text{REPORT}], \text{Instruction}]$$

- **Autoregressive Text Generation:** GPT-2 predicts next token

$$y_t = \arg \max P(y_t | T, y_{<t})$$

Methodology

Our approach uses the [PeFoMed framework](#), a parameter-efficient method for fine-tuning multimodal large language models for medical report generation.

Dataset Preparation

- Base dataset: CheXpert (14 medical condition labels).
- Each image-label pair was used as input to LLaMA to generate descriptive text reports.
- Verified 100 randomly selected samples manually — the generated reports were accurate, label-consistent, and hallucination-free. So proceeded with generating for around 13k samples

Data Split

- Total samples: 10868(after cleaning)
- Split into Train: 9761, Validation: 907, Test: 200

Model Training

- Input: X-ray image
- Output: LLM-generated text report
- Trained for 3 epochs using the custom dataset.

Experimental Design

We conduct a comparative study to evaluate the impact of large-scale, multigranular pre-training on medical report generation quality.

Implementation Details

- Frameworks: PyTorch, HuggingFace Transformers
- Hardware Apple M-series GPU (MPS backend)
- Training Duration \~ 30 mins per epoch processed efficiently
- LLM Used LLaMA for text generation
- PeFoMed Checkpoint Loaded pretrained weights and fine-tuned for 3 epochs

Evaluation Metrics:

- **ROUGE-L** — measures textual overlap between generated and reference reports.
- **METEOR** — evaluates precision, recall, and alignment of phrases.
- **CIDEr** — measures consensus between generated reports and human-written reports.
- **Exact Match (EM)** — checks if generated reports exactly match reference reports.

Results

Training Logs

Epoch 1: Avg Loss: 0.9069

Validation Metrics: ROUGE-L 0.1447, METEOR 0.1146, CIDEr 0.0966, EM 0.0000

Epoch 2: Avg Loss: 0.4718

Validation Metrics: ROUGE-L 0.1589, METEOR 0.1042, CIDEr 0.1003, EM 0.0000

Epoch 3: Avg Loss: 0.4718

Validation Metrics: ROUGE-L 0.1712, METEOR 0.1180, CIDEr 0.1125, EM 0.0000

- **Report Completeness:** More comprehensive coverage of all clinically relevant findings.

Evaluation Results

ROUGE-L 0.129

METEOR 0.095

CIDEr 0.173

EM 0.000

Observations

Across three epochs, the model showed consistent improvements in ROUGE-L, METEOR, and CIDEr. This indicates that the multimodal alignment and LoRA fine-tuning strategy successfully learned image-conditioned report generation.

Comparison with Prior Medical Report Generation Models

Method	Type	METEOR ↑	ROUGE-L ↑	CIDEr ↑
R2Gen [3]	Non-LLMs	0.211	0.377	0.438
BiomedGPT [54]	LLMs	0.146	0.302	0.360
PeFoMed (paper)	LLMs	0.157	0.286	0.462
PeFoMed-Light (Ours)	LLMs	0.095	0.129	0.173

Method	Type	METEOR ↑	ROUGE-L ↑	CIDEr ↑
R2Gen [3]	Non-LLMs	0.211	0.377	0.438
BiomedGPT [54]	LLMs	0.146	0.302	0.360
PeFoMed (paper)	LLMs	0.157	0.286	0.462
PeFoMed-Light (Ours)	LLMs	0.095	0.129	0.173

Conclusion and Future work

The LLM-generated reports from LLaMA proved highly reliable and aligned with the 14 CheXpert labels. PeFoMed successfully learned to map X-ray images to meaningful clinical text.

Future work includes:

- Expanding dataset size
- Training on llama as decoder (if resource provided)
- Replacing vision encoder with Vmamba
- Adding MoE on text decoder and vision encoder part

References

- ❖ **JinlHe, PeFoMed:** Medical Report Generation Framework
- ❖ **GitHub:** [<https://github.com/jinlHe/PeFoMed>]
- ❖ **CheXpert Dataset** — Stanford ML Group
- ❖ **Llama LLM** — Text generation for medical data augmentation
- ❖ **GPT-2 — Text decoder,** [<https://GPT-2-v2.github.io>]
- ❖ **CLIP ViT-L/14 — Vision encoder,** [<https://arxiv.org/pdf/2103.00020>]